



DataCamp DataCamp

Open in app ↗

Sign up

Sign In



Bruce Yang [Follow](#)

Oct 26, 2021 · 7 min read · ✨ · ⏴ Listen

Save



A Data Scientist's Approach for Algorithmic Trading using Deep Reinforcement Learning: An End-to-end Tutorial for Paper Trading

End-to-end tutorial in one Jupyter Notebook: get the data, perform data engineering, train and deploy the DRL agent, place trades, check the portfolio performance

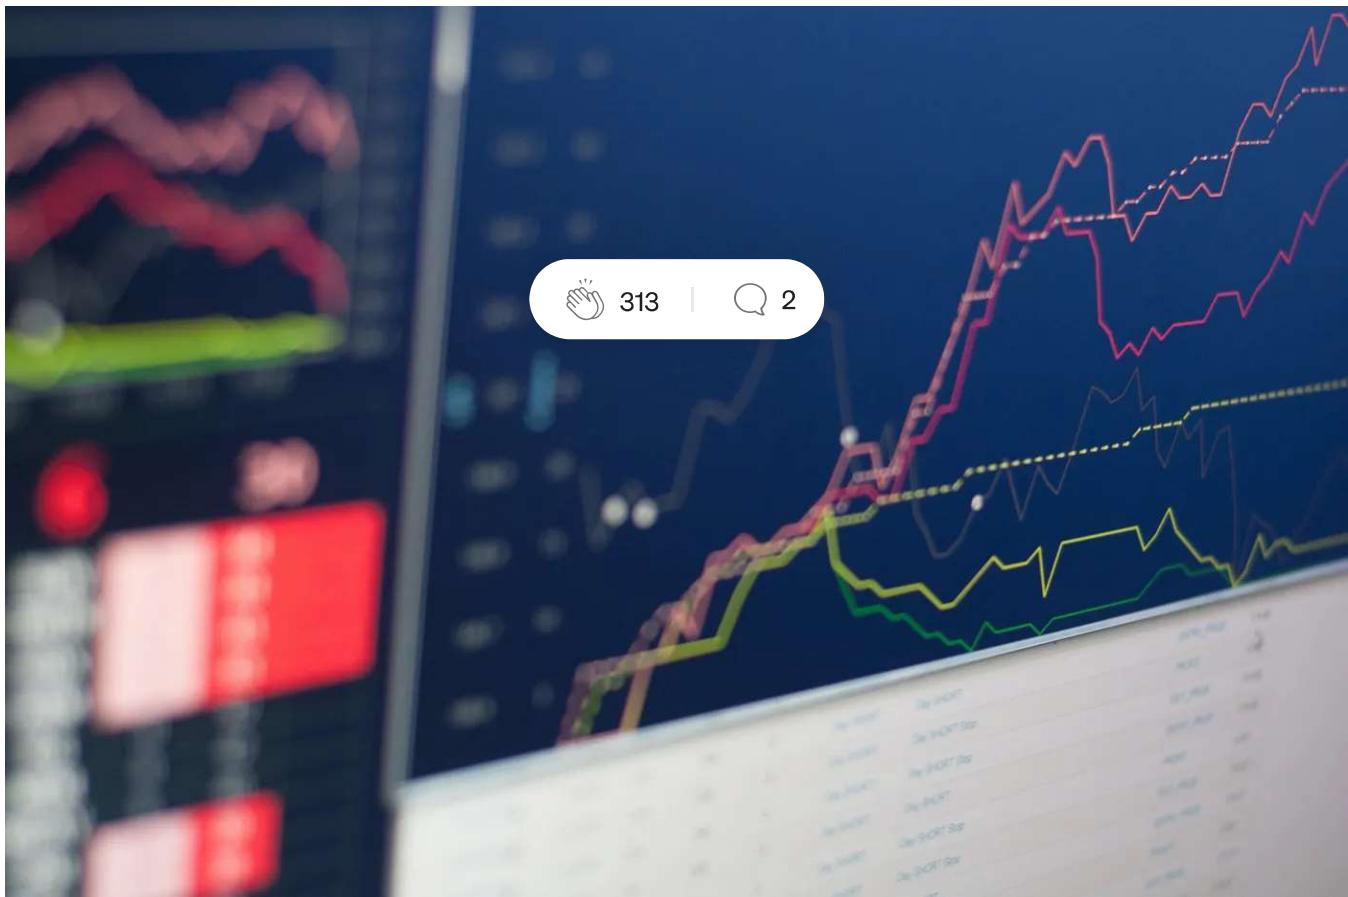


Image by [bash_profile](#) on [Unsplash](#)

Disclaimer: This is NOT a financial advice, and NOT a recommendation to trade real money. Many broker platforms exist for simulated trading (paper trading) which can be used for building and developing the methods discussed. Please use common sense and always first consult a financial advisor before trading or investing.

This article is prepared by Bruce Yang, Jingyang Rui, and [Xiao-Yang Liu](#).

Paper trading is a MUST for algorithmic trading especially for machine learning trading strategies. Most fund would test and fine-tune their strategies through paper trading before shooting real trades.

In quantitative finance, stock trading is essentially making dynamic decisions, namely to decide *where to trade, at what price, and what quantity* over a highly stochastic and complex stock market. As a result, Deep Reinforcement Learning (DRL) provides useful toolkits for stock trading. Taking many complex financial factors into account, DRL trading agents build a multi-factor model and provide algorithmic trading strategies, which are difficult for human traders.

In this article, we use a single Jupyter notebook to show all the steps for paper trading to test our DRL strategy in real life!

After reading this article you will be able to

- Train a DRL agent on *minute* level data of Dow 30
- Deploy the DRL agent to Alpaca trading API using FinRL
- Place trades on Alpaca using virtual money through Jupyter Notebook

What is Paper Trading?

Paper trading is simulated trading that allows people to 1) practice trading securities; 2) test and fine-tune new investment strategies before apply it in live account; 3) teach beginners how to place orders and learn the trading basics.

However, paper trading does not 1) reflect true emotions in live trading; 2) deal with risk aversion; 3) represent real spreads because it's virtual money not real money.

Part 1: Setup Alpaca Paper Trading API



We select Alpaca Trading API mostly because it's free and beginner friendly.

1.1 Key Advantages:

- Commission-free like Robinhood
- Easy to use, good interface
- Good API for algo-trading, beginner friendly
- Unlimited testing of the strategies and bots for free
- No need to be a US resident

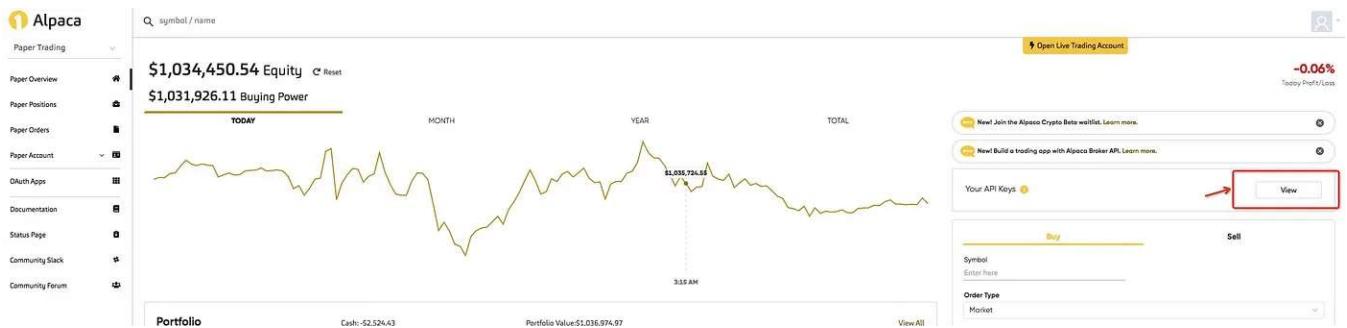
The free rate limit is 200 requests every minute per API key.

1.2 API Endpoints Setup

In [Alpaca's documentation](#), you'll find all the API calls to submit trades, check the portfolio, and utilizing Python. Here are some [good blogs](#) that have very detailed introduction.

To get started in this article, you only need to 1) register an account, 2) generate and keep your **API Key ID & Secret ID** (need to regenerate key if forget).

The screenshot shows the Alpaca website's registration page. On the left, there is a sidebar with links: Live Trading, Paper Accounts (which is highlighted with a red box and a red arrow pointing to it), Documentation, Status Page, Community Slack, and Community Forum. The main content area has a yellow header "Welcome to Alpaca!". Below it, there is a button "Get your live trading account" and two options: "Individual Account" (highlighted in yellow) and "Business Trading Account". Further down, there are links for "Need help? Contact Support!", "Browse our Docs for use cases, sample code, and how to get started.", and "Try your paper trading API first!". At the bottom of the page, there are three sections: "Commission free trading" (with a \$0 icon), "Fractional shares" (with a dollar sign icon), and "Crypto 24/7" (with a Bitcoin icon). Each section has a brief description and a link to more information.



Copyright by AI4Finance-Foundation

Your API Keys i

Endpoint

API Key ID

Secret Key

Your secret key will disappear after refreshing or navigating away from this page.

Copyright by AI4Finance-Foundation

Part 2: Get our Deep Reinforcement Learning Agent Ready!



Image by [eiskonen](#) on [Unsplash](#)

2.0 Install FinRL

FinRL is an open-source framework to help practitioners establish the development pipeline of trading strategies based on **deep reinforcement learning (DRL)**. The FinRL framework allows users to plug in and play with **standard DRL algorithms**. Please refer to [this blog](#) for FinRL installation. This blog uses Google Colab to run Jupyter notebook, please find our code here:

[FinRL/FinRL_PaperTrading_Demo.ipynb at master · AI4Finance-Foundation/FinRL](#)

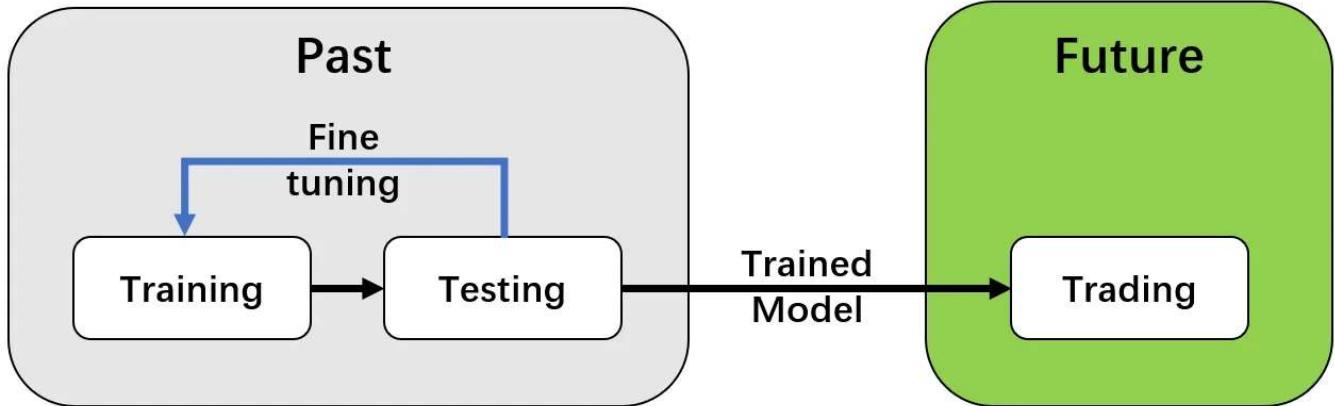
FinRL: Financial Reinforcement Learning. 🔥. Contribute to AI4Finance-Foundation/FinRL development by creating an...

github.com

The “**training-testing**” workflow used by conventional machine learning methods falls short for financial tasks. It splits the data into training set and testing set. On the training data, users select features and tune parameters (k-folds); then make inference on the testing data. However, financial tasks will experience a simulation-to-reality gap between the testing performance and real-live market performance.

Because the testing here is offline backtesting, while the users' goal is to place orders in a real-world market.

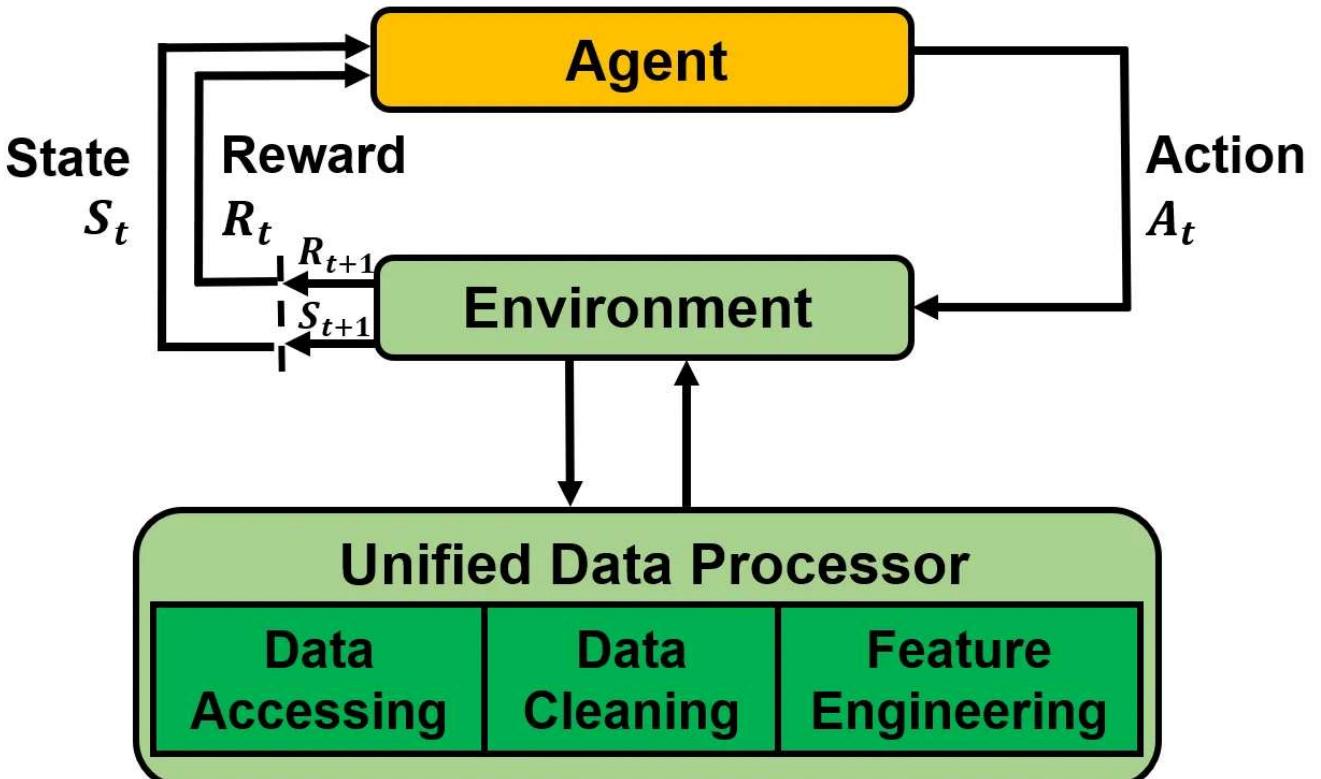
FinRL employs a “training-testing-trading” pipeline to reduce the simulation-to-reality gap.



Copyright by AI4Finance-Foundation

We use historical data (time series) for the “training-testing” part, which is the same as conventional machine learning tasks, and this testing period is for backtesting purpose. For the “trading” part, we use live trading APIs, in this case is Alpaca, allowing users carry out trades directly in a trading system.

2.2 Data Engineering



We establish a standard pipeline for **financial data engineering** in RL, ensuring data of different formats from different sources can be incorporated in a unified framework. Then, we automate this pipeline with a **data processor**, which can *access data, clean data, and extract features* from various data sources with high quality and efficiency. Our data layer provides agility to model deployment.

Step 1: Pick a data source

```
DP = DataProcessor(data_source = 'alpaca',
                    API_KEY = API_KEY,
                    API_SECRET = API_SECRET,
                    APCA_API_BASE_URL = APCA_API_BASE_URL)
```

Step 2: Get the stock ticker list, set start/end date, and specify the time_interval

```
data = DP.download_data(start_date = '2021-10-01',
                        end_date = '2021-10-05',
                        ticker_list = ticker_list,
                        time_interval= '1Min')
```

Step 3: Data Cleaning & Feature Engineering

```
data = DP.clean_data(data)
data = DP.add_technical_indicator(data, TECHNICAL_INDICATORS_LIST)
data = DP.add_vix(data)
```

Step 4: Transform to numpy array format

```
price_array, tech_array, turbulence_array = DP.df_to_array(data,
if_vix='True')
```

	time	open	high	low	close	volume	tic	macd	boll_ub	boll_lb	rsi_30	cci_30	dx_30	close_30_sma	close_60_sma	VIXY
35080	2021-10-05 15:59:00-04:00	143.140	143.180	143.100	143.160	4180.0	IBM	-0.145148	143.866517	143.059983	30.924818	-160.386451	48.082544	143.534000	143.722417	21.68
35081	2021-10-05 15:59:00-04:00	53.955	53.975	53.935	53.945	6274.0	INTC	-0.038827	54.166900	53.925100	29.721071	-169.470945	57.246117	54.060000	54.096000	21.68
35082	2021-10-05 15:59:00-04:00	159.720	159.730	159.540	159.580	4075.0	JNJ	-0.073306	160.073831	159.657669	34.877315	-306.971514	34.150919	159.888667	160.026250	21.68
35083	2021-10-05 15:59:00-04:00	168.870	168.880	168.650	168.650	7616.0	JPM	-0.140191	169.514417	168.669583	33.354857	-173.244738	42.521277	169.178000	169.313083	21.68
35084	2021-10-05 15:59:00-04:00	53.085	53.090	53.055	53.080	15679.0	KO	-0.012753	53.198559	53.068441	42.715325	-190.357440	29.750107	53.138333	53.137750	21.68
35085	2021-10-05 15:59:00-04:00	245.330	245.330	245.180	245.180	2704.0	MCD	-0.269918	246.739285	245.066715	21.785286	-160.170797	79.728448	246.058333	246.202083	21.68
35086	2021-10-05 15:59:00-04:00	177.960	177.960	177.790	177.790	1754.0	MMM	-0.076775	178.495952	177.865548	34.300943	-232.273583	44.417489	178.213500	178.172583	21.68
35087	2021-10-05 15:59:00-04:00	81.595	81.610	81.575	81.590	8480.0	MRK	-0.054943	81.956850	81.510150	38.831341	-142.857143	25.207846	81.732167	81.815083	21.68
35088	2021-10-05 15:59:00-04:00	288.850	288.880	288.530	288.730	24244.0	MSFT	-0.146467	289.571293	288.706207	39.004141	-188.379205	42.062003	289.225667	289.492667	21.68
35089	2021-10-05 15:59:00-04:00	149.580	149.590	149.465	149.520	6883.0	NKE	-0.111984	150.142632	149.492368	37.775018	-166.219839	60.708595	149.912500	149.907083	21.68
35090	2021-10-05 15:59:00-04:00	42.345	42.355	42.290	42.300	19989.0	PFE	-0.027733	42.429359	42.284641	38.597015	-120.575885	23.939389	42.380000	42.431583	21.68
35091	2021-10-05 15:59:00-04:00	139.430	139.430	139.310	139.310	3316.0	PG	-0.083896	139.923238	139.396262	31.750025	-264.842747	48.807575	139.713667	139.735833	21.68
35092	2021-10-05 15:59:00-04:00	88.780	88.780	88.690	88.740	4767.0	RTX	-0.074596	89.244555	88.704945	34.390595	-186.968281	63.034810	89.004667	89.026250	21.68
35093	2021-10-05 15:59:00-04:00	153.570	153.650	153.570	153.650	1320.0	TRV	-0.038149	153.850881	153.514619	46.886586	-74.290919	24.670390	153.720500	153.740417	21.68
35094	2021-10-05 15:59:00-04:00	393.900	393.900	393.200	393.200	2778.0	UNH	-0.337330	395.924222	393.719278	35.160318	-314.024950	44.275058	394.942833	395.146833	21.68
35095	2021-10-05 15:59:00-04:00	224.340	224.380	224.200	224.270	5364.0	V	-0.232614	225.621581	224.186419	27.798707	-173.469388	71.642315	225.030167	225.164333	21.68
35096	2021-10-05 15:59:00-04:00	54.550	54.555	54.500	54.500	13804.0	VZ	-0.026883	54.660802	54.528688	29.257106	-236.286920	59.868748	54.610167	54.652333	21.68
35097	2021-10-05 15:59:00-04:00	47.100	47.100	47.080	47.090	3300.0	WBA	-0.000467	47.216146	47.085354	47.064782	-62.975239	26.794769	47.125000	47.107500	21.68
35098	2021-10-05 15:59:00-04:00	136.690	136.710	136.600	136.620	5695.0	WMT	-0.076694	137.037700	136.663800	31.557994	-218.758646	61.392805	136.904833	136.971083	21.68
35099	2021-10-05 15:59:00-04:00	61.655	61.670	61.590	61.610	24331.0	XOM	-0.060220	61.912768	61.574732	38.426421	-132.690302	60.013011	61.791167	61.855167	21.68

Copyright by AI4Finance-Foundation

2.3 Train our DRL agent

We only need to provide some basic parameters and model hyperparameters to the train() function, after training finished, it will output the trained model to the folder we specified (cwd) along with the learning rate plot.

```
#demo for elegantrl
ERL_PARAMS = {"learning_rate": 3e-5,
               "batch_size": 2048,
               "gamma": 0.99,
               "seed": 312,
               "net_dimension": 512}

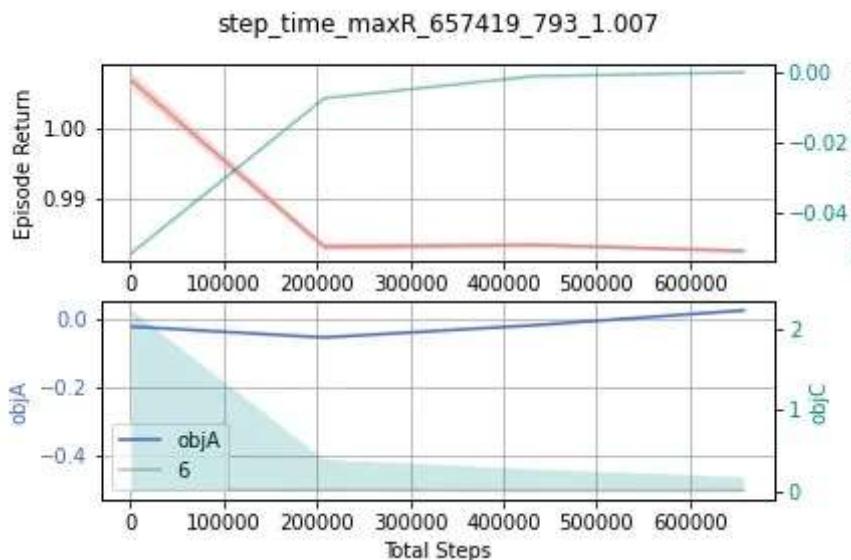
train(start_date = '2021-10-11',
       end_date = '2021-10-15',
       ticker_list = ticker_list,
       data_source = 'alpaca',
       time_interval= '1Min',
       technical_indicator_list= TECHNICAL_INDICATORS_LIST,
       drl_lib='elegantrl',
       env=env,
       model_name='ppo',
       API_KEY = API_KEY,
       API_SECRET = API_SECRET,
       APCA_API_BASE_URL = APCA_API_BASE_URL,
       erl_params=ERL_PARAMS,
       cwd='./papertrading_erl',
       total_timesteps=1e5)
```

```

Alpaca successfully connected
Data before 2021-10-11T15:59:00-04:00 is successfully fetched
Data before 2021-10-12T15:59:00-04:00 is successfully fetched
Data before 2021-10-13T15:59:00-04:00 is successfully fetched
Data before 2021-10-14T15:59:00-04:00 is successfully fetched
Data before 2021-10-15T15:59:00-04:00 is successfully fetched
Data clean finished!
Successfully add technical indicators
Data before 2021-10-11T15:59:00-04:00 is successfully fetched
Data before 2021-10-12T15:59:00-04:00 is successfully fetched
Data before 2021-10-13T15:59:00-04:00 is successfully fetched
Data before 2021-10-14T15:59:00-04:00 is successfully fetched
Data before 2021-10-15T15:59:00-04:00 is successfully fetched
Data clean finished!
Successfully transformed into array
| Remove cwd: ./papertrading_erl
#####
ID      Step    maxR |      avgR   stdR   avgS   stdS |      expR     objC     etc.
0  3.90e+03  1.01 |      1.01   0.0   1948      0 |     -0.05   1.96   0.04  -0.50
0  3.90e+03  1.01 |      1.01   0.0   1948      0 |     -0.05   1.96   0.04  -0.50
| UsedTime:      149 | SavedDir: ./papertrading_erl

```

Copyright by AI4Finance-Foundation



Copyright by AI4Finance-Foundation

2.4 Test our DRL agent

The purpose of backtesting is to tune model hyperparameters, make sure that we get a positive and valid return/reward (model converges), and get a good performing strategy.

```

#demo for elegantrl

account_value_erl=
test(start_date = '2021-10-18',
      end_date = '2021-10-19',
      ticker_list = ticker_list,

```

```

data_source = 'alpaca',
time_interval= '1Min',
technical_indicator_list= TECHNICAL_INDICATORS_LIST,
drl_lib='elegantrl',
env=env,
model_name='ppo',
API_KEY = API_KEY,
API_SECRET = API_SECRET,
APCA_API_BASE_URL = APCA_API_BASE_URL,
cwd='./papertrading_erl',
net_dimension = 512)

```

```

Alpaca successfully connected
Data before 2021-10-18T15:59:00-04:00 is successfully fetched
Data before 2021-10-19T15:59:00-04:00 is successfully fetched
Data clean finished!
Successfully add technical indicators
Data before 2021-10-18T15:59:00-04:00 is successfully fetched
Data before 2021-10-19T15:59:00-04:00 is successfully fetched
Data clean finished!
Successfully transformed into array
price_array: 780
Test Finished!
episode_return 1.0026875870882481

```

Copyright by AI4Finance-Foundation

2.5 Use full data to train

After the hyperparameters is fixed, we need to use all data available up to the point to train the model because model needs to dynamically adapt to new patterns in data, in case of [concept drift](#).

```

#demo for elegantrl
ERL_PARAMS = {"learning_rate": 3e-5,
               "batch_size": 2048,
               "gamma": 0.99,
               "seed":312,
               "net_dimension":512}

train(start_date = '2021-10-11',
      end_date = '2021-10-19',
      ticker_list = ticker_list,
      data_source = 'alpaca',
      time_interval= '1Min',
      technical_indicator_list= TECHNICAL_INDICATORS_LIST,
      drl_lib='elegantrl',
      env=env,
      model_name='ppo',
      API_KEY = API_KEY,
      API_SECRET = API_SECRET,

```

```
APCA_API_BASE_URL = APCA_API_BASE_URL,  
erl_params=ERL_PARAMS,  
cwd='./papertrading_erl',  
total_timesteps=1e5)
```

Part 3: Deploy our DRL Agent to Alpaca Paper Trading API

3.1 Deploy our agent from the saved file

We load the trained DRL model, and connect it to AlpacaPaperTrading environment to start shooting trades.

```
#demo for elegantrl  
  
paper_trading_erl =  
AlpacaPaperTrading(ticker_list = DOW_30_TICKER,  
time_interval = '1Min',  
drl_lib = 'elegantrl',  
agent = 'ppo',  
cwd = './papertrading_erl',  
net_dim = 512,  
state_dim = state_dim,  
action_dim= action_dim,  
API_KEY = API_KEY,  
API_SECRET = API_SECRET,  
APCA_API_BASE_URL = APCA_API_BASE_URL,  
tech_indicator_list = TECHNICAL_INDICATORS_LIST,  
turbulence_thresh=30,  
max_stock=1e2)  
  
paper_trading_erl.run()
```

```
• Waiting for market to open...
Market opened.
Successfully add technical indicators
Successfully transformed into array
30
Quantity is 0, order of | 0 AMGN sell | not completed.
Quantity is 0, order of | 0 HON sell | not completed.
Quantity is 0, order of | 0 IBM sell | not completed.
Quantity is 0, order of | 0 PG sell | not completed.
Quantity is 0, order of | 0 TRV sell | not completed.
Quantity is 0, order of | 0 UNH sell | not completed.
Quantity is 0, order of | 0 WBA sell | not completed.
Quantity is 0, order of | 0 WMT sell | not completed.
Quantity is 0, order of | 0 DIS sell | not completed.
Market order of | 20 AXP buy | completed.
Market order of | 33 AAPL buy | completed.
Market order of | 11 CAT buy | completed.
Market order of | 50 CSCO buy | completed.
Market order of | 33 HD buy | completed.
Market order of | 32 INTC buy | completed.
Market order of | 27 JNJ buy | completed.
Market order of | 26 MMM buy | completed.
Market order of | 62 MSFT buy | completed.
Market order of | 27 NKE buy | completed.
Market order of | 44 CRM buy | completed.
Market order of | 28 V buy | completed.
```

Copyright by AI4Finance-Foundation

3.2 Alpaca Paper Trading Environment for Reinforcement Learning

Environment is the key for reinforcement learning, FinRL provides the [Alpaca Paper Trading Env](#), it connects the DRL agent to the Alpaca Trading API and automatically place trades.

1. Load agent
2. Connect to Alpaca trading API
3. Read trading time interval
4. Read trading settings
5. Initialize account

6. Wait for market to open.
7. Get DRL states.
8. Submit order to place trades.

```
1 import datetime
2 import threading
3 from finrl.neo_finrl.data_processors.processor_alpaca import AlpacaProcessor
4 from elegantrl.run import *
5 import alpaca_trade_api as tradeapi
6 import time
7 import pandas as pd
8 import numpy as np
9 import torch
10 import sys
11 import os
12 import gym
13
14 class AlpacaPaperTrading():
15
16     def __init__(self,ticker_list, time_interval, drl_lib, agent, cwd, net_dim,
17                  state_dim, action_dim, API_KEY, API_SECRET,
18                  APCA_API_BASE_URL, tech_indicator_list, turbulence_thresh=30,
19                  max_stock=1e2, latency = None):
20         #load agent
21         self.drl_lib = drl_lib
22         if agent =='ppo':
23             if drl_lib == 'elegantrl':
24                 from elegantrl.agent import AgentPPO
25                 #load agent
26                 try:
27                     agent = AgentPPO()
28                     agent.init(net_dim, state_dim, action_dim)
29                     agent.save_or_load_agent(cwd=cwd, if_save=False)
30                     self.act = agent.act
31                     self.device = agent.device
32                 except:
33                     raise ValueError('Fail to load agent!')
34
35         elif drl_lib == 'rllib':
36             from ray.rllib.agents import ppo
37             from ray.rllib.agents.ppo.ppo import PPOTrainer
38
39             config = ppo.DEFAULT_CONFIG.copy()
40             config['env'] = StockEnvEmpty
41             config["log_level"] = "WARN"
42             config['env_config'] = {'state_dim':state_dim,
43                                    'action_dim':action_dim,}
44             trainer = PPOTrainer(env=StockEnvEmpty, config=config)
45             trainer.restore(cwd)
46             try:
47                 trainer.restore(cwd)
48                 self.agent = trainer
49             except:
50                 self.agent = None
```

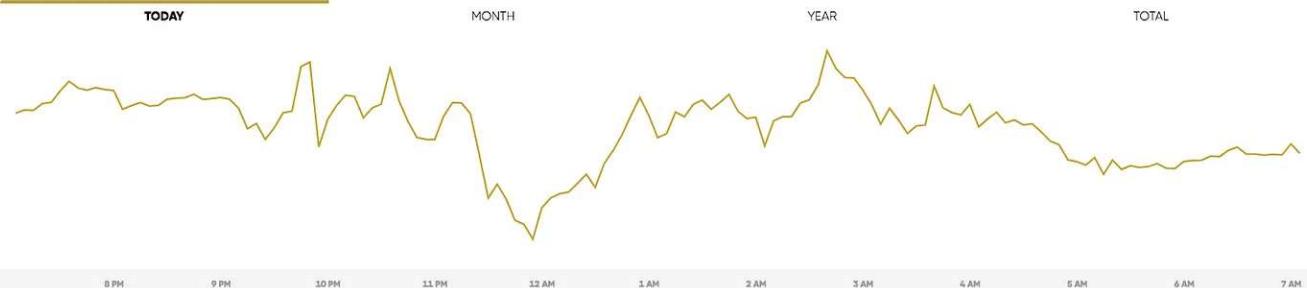
```

40             self._agent = Trainer()
41
42         print("Restoring from checkpoint path", cwd)
43
44     except:
45
46         raise ValueError('Fail to load agent!')
47
48
49     elif drl_lib == 'stable baselines3':

```

\$1,034,450.54 Equity C Reset

\$1,031,926.11 Buying Power



Portfolio	Cash: -\$2,524.43	Portfolio Value: \$1,036,974.97	View All
Asset	Price	Quantity	Market Value
V +0.00%	\$231.23	252	\$58,269.96
NKE -0.17%	\$163.20	453	\$73,929.60
MSFT -0.05%	\$309.00	1006	\$310,854.00
MMM +0.00%	\$180.80	299	\$54,059.20
JNJ +0.00%	\$163.72	254	\$41,584.88
INTC -0.08%	\$49.42	263	\$12,997.46
...			

/ ± CONNECT TO ALPACA TRADING API

```

72     try:
73
74         self.alpaca = tradeapi.REST(API_KEY,API_SECRET,APCA_API_BASE_URL, 'v2')
75
76     except:

```

Positions



Liquidate Selected Positions										
Stock	Shares	Price	Avg Entry	Cost Basis	Market Value	Today's P/L (%)	Today's P/L (\$)	Total P/L (\$)	Total P/L (%)	
V	252	\$231.23	\$224.44	\$56,558.24	\$58,269.96	0.00%	\$0.00	+\$3,711.72	+\$3,711.72	
NKE	453	\$163.20	\$157.18	\$71,203.16	\$73,929.60	-0.17%	-\$126.84	+\$2,726.44	+\$2,726.44	
MSFT	1006	\$309.00	\$300.02	\$301,822.36	\$310,854.00	-0.05%	-\$106.96	+\$9,031.64	+\$9,031.64	
MMM	299	\$180.80	\$178.42	\$53,348.20	\$54,059.20	0.00%	\$549.20	+\$711.00	+\$711.00	
JNJ	254	\$163.72	\$160.39	\$40,739.56	\$41,584.88	0.00%	\$845.32	+\$844.92	+\$844.92	
INTC	263	\$49.42	\$55.01	\$13,941.65	\$12,997.46	-0.08%	-\$944.19			
HD	442	\$365.31	\$341.54	\$157,792.10	\$168,773.22	-0.19%	-\$114.16	+\$10,981.12	+\$10,981.12	
CSCO	820	\$55.01	\$55.33	\$45,307.24	\$45,508.20	-0.18%	-\$80.96	-0.58%	-0.58%	
CRM	389	\$192.49	\$287.23	\$111,731.50	\$113,778.61	-0.02%	-\$27.23	+1.03%	+1.03%	
CAT	44	\$200.63	\$199.14	\$8,027.72	\$8,027.72	-0.01%	-\$0.80	+0.00%	+0.00%	

80 self._time_interval = 60 * 5

```

86     elif time_interval == '15Min':
87
88         self.time_interval = 60 * 15
89
90
91     #read trading settings
92     self.tech_indicator_list = tech_indicator_list
93     self.turbulence_thresh = turbulence_thresh
94     self.max_stock = max_stock
95

```

Stock Trading (Paper trading)



Year 2021

```
116     print('latency for data processing: ', latency)
117     return latency
118
119     def run(self):
120         orders = self.alpaca.list_orders(status="open")
121         for order in orders:
122             self.alpaca.cancel_order(order.id)
123
124         # Wait for market to open.
125         print("Waiting for market to open...")
126         tAMO = threading.Thread(target=self.awaitMarketOpen)
127         tAMO.start()
128         tAMO.join()
129         print("Market opened.")
130         while True:
131
132             # Figure out when the market will close so we can prepare to sell beforehand.
133             clock = self.alpaca.get_clock()
134             closingTime = clock.next_close.replace(tzinfo=datetime.timezone.utc).timestamp()
135             currTime = clock.timestamp.replace(tzinfo=datetime.timezone.utc).timestamp()
136             self.timeToClose = closingTime - currTime
137
138             if(self.timeToClose < (60)):
139                 # Close all positions when 1 minutes til market close.
140                 print("Market closing soon. Stop trading.")
141                 break
142
143             '''# Close all positions when 1 minutes til market close
```

```
144     print("Market closing soon. Closing positions.")
145
146     positions = self.alpaca.list_positions()
147     for position in positions:
148         if(position.side == 'long'):
149             orderSide = 'sell'
150         else:
151             orderSide = 'buy'
152             qty = abs(int(float(position.qty)))
153             respSO = []
154             tSubmitOrder = threading.Thread(target=self.submitOrder(qty, position.symbol, or
155             tSubmitOrder.start()
156             tSubmitOrder.join()
157
158             # Run script again after market close for next trading day.
159             print("Sleeping until market close (15 minutes).")
160             time.sleep(60 * 15)'''
161
162         else:
163             trade = threading.Thread(target=self.trade)
164             trade.start()
165             trade.join()
166             last_equity = float(self.alpaca.get_account().last_equity)
167             cur_time = time.time()
168             self.equities.append([cur_time, last_equity])
169             time.sleep(self.time_interval)
170
171     def awaitMarketOpen(self):
172         isOpen = self.alpaca.get_clock().is_open
173         while(not isOpen):
174             clock = self.alpaca.get_clock()
175             openingTime = clock.next_open.replace(tzinfo=datetime.timezone.utc).timestamp()
176             currTime = clock.timestamp.replace(tzinfo=datetime.timezone.utc).timestamp()
177             timeToOpen = int((openingTime - currTime) / 60)
178             print(str(timeToOpen) + " minutes til market open.")
179             time.sleep(60)
180             isOpen = self.alpaca.get_clock().is_open
181
182     def trade(self):
183         state = self.get_state()
184
185         if self.drl_lib == 'elegantrl':
186             with torch.no_grad():
187                 s_tensor = torch.as_tensor((state,), device=self.device)
188                 a_tensor = self.act(s_tensor)
189                 action = a_tensor.detach().cpu().numpy()[0]
190
```

```

191     action = (action * self.max_stock).astype(int)
192
193     elif self.drl_lib == 'rllib':
194         action = self.agent.compute_single_action(state)
195
196     elif self.drl_lib == 'stable_baselines3':
197         action = self.model.predict(state)[0]
198
199 else:
200     raise ValueError('The DRL library input is NOT supported yet. Please check your ir'
201
202     self.stocks_cd += 1
203     if self.turbulence_bool == 0:
204         min_action = 10 # stock_cd
205         for index in np.where(action < -min_action)[0]: # sell_index:
206             sell_num_shares = min(self.stocks[index], -action[index])
207             qty = abs(int(sell_num_shares))
208             respSO = []
209             tSubmitOrder = threading.Thread(target=self.submitOrder(qty, self.stockUnivers
210             tSubmitOrder.start()
211             tSubmitOrder.join()
212             self.cash = float(self.alpaca.get_account().cash)
213             self.stocks_cd[index] = 0
214
215             for index in np.where(action > min_action)[0]: # buy_index:
216                 if self.cash < 0:
217                     tmp_cash = 0
218                 else:
219                     tmp_cash = self.cash
220                     buy_num_shares = min(tmp_cash // self.price[index], abs(int(action[index])))
221                     qty = abs(int(buy_num_shares))
222                     respSO = []
223                     tSubmitOrder = threading.Thread(target=self.submitOrder(qty, self.stockUnivers
224                     tSubmitOrder.start()
225                     tSubmitOrder.join()
226                     self.cash = float(self.alpaca.get_account().cash)
227                     self.stocks_cd[index] = 0
228
229             else: # sell all when turbulence
230                 positions = self.alpaca.list_positions()
231                 for position in positions:
232                     if(position.side == 'long'):
233                         orderSide = 'sell'
234                     else:
235                         orderSide = 'buy'
236                     qty = abs(int(float(position.qty)))
237                     respSO = []
238                     tSubmitOrder = threading.Thread(target=self.submitOrder(atv. position.symbol,

```

```
239         tSubmitOrder.start()
240         tSubmitOrder.join()
241
242         self.stocks_cd[:] = 0
243
244
245     def get_state(self):
246         alpaca = AlpacaProcessor(api=self.alpaca)
247         price, tech, turbulence = alpaca.fetch_latest_data(ticker_list = self.stockUniverse, t
248                                         tech_indicator_list=self.tech_indicator_]
249         turbulence_bool = 1 if turbulence >= self.turbulence_thresh else 0
250
251         turbulence = (self.sigmoid_sign(turbulence, self.turbulence_thresh) * 2 ** -5).astype(
252
253         tech = tech * 2 ** -7
254         positions = self.alpaca.list_positions()
255         stocks = [0] * len(self.stockUniverse)
256         for position in positions:
257             ind = self.stockUniverse.index(position.symbol)
258             stocks[ind] = (abs(int(float(position.qty))))))
259
260         stocks = np.asarray(stocks, dtype = float)
261         cash = float(self.alpaca.get_account().cash)
262         self.cash = cash
263         self.stocks = stocks
264         self.turbulence_bool = turbulence_bool
265         self.price = price
266
267
268
269         amount = np.array(max(self.cash, 1e4) * (2 ** -12), dtype=np.float32)
270         scale = np.array(2 ** -6, dtype=np.float32)
271         state = np.hstack((amount,
272                           turbulence,
273                           self.turbulence_bool,
274                           price * scale,
275                           self.stocks * scale,
276                           self.stocks_cd,
277                           tech,
278                           )).astype(np.float32)
279         print(len(self.stockUniverse))
280         return state
281
282     def submitOrder(self, qty, stock, side, resp):
283         if(qty > 0):
284             try:
285                 self.alpaca.submit_order(stock, qty, side, "market", "day")
```

```
286         print("Market order of | " + str(qty) + " " + stock + " " + side + " | completed.")
Reinforcement Learning    'e'  Stock Trading      Data Science      Feature Engineering
287         except:
288             print("Order of | " + str(qty) + " " + stock + " " + side + " | did not go through")
289             resp.append(False)
290
291     else:
292         print("Quantity is 0, order of | " + str(qty) + " " + stock + " " + side + " | not com-
293         resp.append(True)
294
295     @staticmethod
296     def sigmoid_sign(ary, thresh):
297         def sigmoid(x):
298             return 1 / (1 + np.exp(-x * np.e)) - 0.5
299
300         return sigmoid(ary / thresh) * thresh
301
302 class StockEnvEmpty(gym.Env):
303     #Empty Env used for loading rllib agent
304     def __init__(self, config):
305         state_dim = config['state_dim']
306         action_dim = config['action_dim']
307         self.observation_space = gym.spaces.Box(low=-3000, high=3000, shape=(state_dim,), dtype=
308         self.action_space = gym.spaces.Box(low=-1, high=1, shape=(action_dim,), dtype=np.float32)
309
310     def reset(self):
311         return
```